

Team Raspberry - Image Classification

Kristian Wahlroos Ilkka Vähämaa Sean Lang

December 15, 2017

Overview

- ① Methods
- ② Parametrization
- ③ Final System
- ④ Results

Methods

Data representation

- Treat every image as a 3D-tensor (RGB)
 - Repeat the value of grayscale images three times
 - Colorized are handled as the original tensors
- Original data has 14 labels, we used 15
 - Extra one for the unclassified images
 - One-hot encoded labels

Methods

Data processing

- Read images in batches of size 2000
 - Helps to avoid filling the RAM
- Normalize the pixel values between $[0.0, 1.0]$
- For every batch augmenting the data
 - Provided by Keras
 - Centerify, shear, zoom, rotate and flip
 - To get more variation and samples from classes with few labels

Methods

Class weights 1/2

- Classes are very unbalanced




Figure: Class distribution

Methods

Class weights 2/2

- We tackled this problem by custom weights per class
 - Giving them at training phase

Class weight function

$$S(c_i; \lambda) = \ln \left(\lambda \frac{\sum_c |c|}{|c_i|} \right)$$

$$W(c_i; \lambda) = \max(S(c_i; \lambda), 1)$$

Methods

Network topology

- One network that outputs 15 classes
- Four convolution layers all followed by max pooling
 - Filters 16, 32, 32, 64
 - Kernel size 3×3
 - Max pool size 2×2
 - ReLU as activation function
- After pooling flattening via dropout to dense layer with sigmoid activation
 - Dropout value: 0.4
- Very simple network

Methods

Loss function 1/2

- Categorical crossentropy wouldn't work as one image can be in many classes
- Binary crossentropy was suggested in many forum posts
 - Still not viable solution when there are many overlapping categories
 - Loss is too forgiving for giving 0 labels

Methods

Loss function 2/2

- Solution: "custom" loss function **BP-MLL**^{*}
 - Actually taken directly from the paper [1][†]
 - Designed for multi-label problems
 - Implementation for Keras can be found from internet
 - Punishes more from just giving 0 labels

$$E = \sum_{i=1}^m \frac{1}{|Y_i| |\bar{Y}_i|} \sum_{(k,l) \in Y_i \times \bar{Y}_i} \exp(-(c_k^i - c_l^i))$$

^{*}Backpropagation for Multilabel Learning

[†][1] *Multilabel Neural Networks with Applications to Functional Genomics and Text Categorization*, 2006

Methods

Validation

- Per batch, 10% of the data is randomly selected
- This subset is left out from the training phase
- Validated against in the final step
- With F1-score, we also inspected
 - Binary accuracy
 - Categorical accuracy
 - Hamming loss
 - Micro averaged precision score

Parametrization

Tweaks

- 1 "Default"
- 2 Increased deeply-connected layers
- 3 Adagrad optimizer
- 4 Nadam optimizer
- 5 More convolutions
- 6 Even more convolutions
- 7 Reverse convolution triangle
- 8 Learning Rate Adjustments
 - 1 `lr=0.0005`
 - 2 `lr=0.000333`
 - 3 `lr=0.002`
 - 4 `lr=0.005`
- 9 Activation Functions
 - 1 Leaky ReLU ($\alpha = 0.3$)
 - 2 `tanh`

Parametrization

Results


Model n.	F1	HL [‡]
1	0.462	0.125
2	0.452	0.128
3	0.459	0.125
4	0.469	0.123
5	0.462	0.124
6	0.462	0.126
7	0.463	0.125
8.1	0.464	0.124
8.2	0.457	0.126
8.3	0.465	0.124
8.4	0.457	0.125
9.1	0.083	0.796
9.2	0.378	0.184

[‡]Hamming Loss


Parametrization

Training


(a) “Default” (1)




(b) More Dense Layers (2)



(c) More Convolutional Layers (6)



(d) tanh (9.2)



Final System

Hyper-parameters of the final system

- Same as described in Methods section
- 2 epochs per batch
- **BP-MLL** Loss
- RMSprop optimizer (learning rate of 0.002 [twice default])

Results

Accuracy and Loss




Figure: Final Model (trained over 250 epochs)

- 15 label F1: 0.467
- 14 label F1: 0.499
- Hamming loss: 0.123

Results


Sample Output

0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0

...

Results

Confusion Matrix



- Always predicts same result for each label on every picture