Team Raspberry - Image Classification

Kristian Wahlroos   Ilkka Vähämaa   Sean Lang

December 15th, 2017

Overview

Methods
Data representation

- ▶ Treat every image as a 3D-tensor (RGB)
  - ▶ Repeat the value of grayscale images three times
  - ▶ Colorized are handled as the original tensors
- ▶ Original data has 14 labels, we used 15
  - ▶ Extra one for the unclassified images
  - ▶ One-hot encoded labels

Methods

- ▶ Read images in batches of size 2000
  - ▶ Helps to avoid filling the RAM
- ▶ Normalize the pixel values between $[0.0, 1.0]$
- ▶ For every batch augmenting the data
  - ▶ Provided by Keras
  - ▶ Centerify, shear, zoom, rotate and flip
  - ▶ To get more variation and samples from classes with few labels

Methods

▶ Classes are very unbalanced

Figure: Class distribution

Methods

- We tackled this problem by custom weights per class
  - Giving them at training phase

Class weight function

$$S(c_i; \lambda) = \ln \left(\lambda \frac{\sum_c |c|}{|c_i|} \right)$$

$$W(c_i; \lambda) = \max(S(c_i; \lambda), 1)$$

- One network that outputs 15 classes
- Four convolution layers all followed by max pooling
    - Filters 16, 32, 32, 64
    - Kernel size 3x3
    - Max pool size 2x2
    - ReLU as activation function
- After pooling flattening via dropout to dense layer with sigmoid activation
    - Dropout value: 0.4
- Very simple network

Methods
Loss function 1/2

- ▶ Categorical crossentropy wouldn't work as one image can be in many classes
- ▶ Binary crossentropy was suggested in many forum posts
  - ▶ Still not viable solution when there are many overlapping categories
  - ▶ Loss is too forgiving for giving 0 labels

Methods
Loss function 2/2

- ► Solution: "custom" loss function **BP-MLL**[*]
    - ► Actually taken directly from the paper [1][†]
    - ► Designed for multi-label problems
    - ► Implementation for Keras can be found from internet
    - ► Punishes more from just giving 0 labels

$$E = \sum_{i=1}^{m} \frac{1}{|Y_i||\bar{Y}_i|} \sum_{(k,l) \in Y_i \times \bar{Y}_i} \exp(-(c_k^i - c_l^i))$$

---

[*]Backpropagation for Multilabel Learning

[†][1] *Multilabel Neural Networks with Applications to Functional Genomics and Text Categorization*, 2006

Methods

Validation

- Per batch, 10% of the data is randomly selected
- This subset is left out from the training phase
- Validated against in the final step
- With F1-score, we also inspected
    - Binary accuracy
    - Categorical accuracy
    - Hamming loss
    - Micro averaged precision score

Parametrization results

Table: Parameter Tryouts

Model n.	F1	HL[‡]
1	0.462	0.125
2	0.452	0.128
3	0.459	0.125
4	0.469	0.123
5	0.462	0.124
6	0.462	0.126
7	0.463	0.125
8.1	0.464	0.124
8.2	0.457	0.126
8.3	0.465	0.124
8.4	0.457	0.125
9.1	0.083	0.796
9.2	0.378	0.184

[‡]Hamming Loss

Parametrization

Parameterization Tweaks

1. "Default"
2. Increased deeply-connected layers
3. Adagrad optimizer
4. Nadam optimizer
5. More convolutions
6. Even more convolutions
7. Reverse convolution triangle
8. Learning Rate Adjustments
   - 8.1 `lr=0.0005`
   - 8.2 `lr=0.000333`
   - 8.3 `lr=0.002`
   - 8.4 `lr=0.005`
9. Activation Functions
   - 9.1 Leaky ReLU ($\alpha = 0.3$)
   - 9.2 `tanh`

Final System

Hyper-parameters of the final system

- Same as described in Methods section
- 1 epoch
- BP_MLL_Loss
- RMSprop optimizer (learning rate of 0.002 [twice default])

Results

- Needs some fancy plots here

Results
Sample Output

```
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0
```

. . .

Results

Confusion Matrix

- Always predicts same result for each label on every picture

Results