

Team Raspberry - Image Classification

Kristian Wahlroos Ilkka Vähämaa Sean Lang

December 15th, 2017

Overview

Methods

Parametrization

Final System

Results

Methods

Data representation

- ▶ Treat every image as a 3D-tensor (RGB)
 - ▶ Repeat the value of grayscale images three times
 - ▶ Colorized are handled as the original tensors
- ▶ Original data has 14 labels, we used 15
 - ▶ Extra one for the unclassified images
 - ▶ One-hot encoded labels

Methods

Data processing

- ▶ Read images in batches of size 2000
 - ▶ Helps to avoid filling the RAM
- ▶ Normalize the pixel values between $[0.0, 1.0]$
- ▶ For every batch augmenting the data
 - ▶ Provided by Keras
 - ▶ Centerify, shear, zoom, rotate and flip
 - ▶ To get more variation and samples from classes with few labels

Methods

Class weights 1/2

- Classes are very unbalanced




Figure: Class distribution

Methods

Class weights 2/2

- ▶ We tackled this problem by custom weights per class
 - ▶ Giving them at training phase

Class weight function

$$S(c_i; \lambda) = \ln \left(\lambda \frac{\sum_c |c|}{|c_i|} \right)$$

$$W(c_i; \lambda) = \max(S(c_i; \lambda), 1)$$

Methods

Network topology

- ▶ One network that outputs 15 classes
- ▶ Four convolution layers all followed by max pooling
 - ▶ Filters 16, 32, 32, 64
 - ▶ Kernel size 3x3
 - ▶ Max pool size 2x2
 - ▶ ReLU as activation function
- ▶ After pooling flattening via dropout to dense layer with sigmoid activation
 - ▶ Dropout value: 0.4
- ▶ Very simple network

Methods

Loss function 1/2

- ▶ Categorical crossentropy wouldn't work as one image can be in many classes
- ▶ Binary crossentropy was suggested in many forum posts
 - ▶ Still not viable solution when there are many overlapping categories
 - ▶ Loss is too forgiving for giving 0 labels

Methods

Loss function 2/2

- ▶ Solution: "custom" loss function **BP-MLL**^{*}
 - ▶ Actually taken directly from the paper [1][†]
 - ▶ Designed for multi-label problems
 - ▶ Implementation for Keras can be found from internet
 - ▶ Punishes more from just giving 0 labels

$$E = \sum_{i=1}^m \frac{1}{|Y_i| |\bar{Y}_i|} \sum_{(k,l) \in Y_i \times \bar{Y}_i} \exp(-(c_k^i - c_l^i))$$

^{*}Backpropagation for Multilabel Learning

[†][1] *Multilabel Neural Networks with Applications to Functional Genomics and Text Categorization*, 2006

Methods

Validation

- ▶ Per batch, 10% of the data is randomly selected
- ▶ This subset is left out from the training phase
- ▶ Validated against in the final step
- ▶ With F1-score, we also inspected
 - ▶ Binary accuracy
 - ▶ Categorical accuracy
 - ▶ Hamming loss
 - ▶ Micro averaged precision score

Parametrization

Parameterization Tweaks

1. “Default”
2. Increased deeply-connected layers
3. Adagrad optimizer
4. Nadam optimizer
5. More convolutions
6. Even more convolutions
7. Reverse convolution triangle
8. Learning Rate Adjustments
 - 8.1 `lr=0.0005`
 - 8.2 `lr=0.000333`
 - 8.3 `lr=0.002`
 - 8.4 `lr=0.005`
9. Activation Functions
 - 9.1 Leaky ReLU ($\alpha = 0.3$)
 - 9.2 `tanh`

Parametrization results

Table: Parameter Tryouts

Model n.	F1	HL [‡]
1	0.462	0.125
2	0.452	0.128
3	0.459	0.125
4	0.469	0.123
5	0.462	0.124
6	0.462	0.126
7	0.463	0.125
8.1	0.464	0.124
8.2	0.457	0.126
8.3	0.465	0.124
8.4	0.457	0.125
9.1	0.083	0.796
9.2	0.378	0.184

[‡]Hamming Loss

Final System

Hyper-parameters of the final system

- ▶ Same as described in Methods section
- ▶ 1 epoch
- ▶ BP_MLL_Loss
- ▶ RMSprop optimizer (learning rate of 0.002 [twice default])

Results

- ▶ Needs some fancy plots here

Results


Sample Output

```
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
```

...

Results

Confusion Matrix



- Always predicts same result for each label on every picture

Results




Figure: Soemthing